


# TTIC 31110

## Speech Technologies


April 30, 2020

## Questions from last time: Continuous speech recognition with HMMs


Most basic model: String together word HMMs to make sentence HMM, using a grammar (Note: Start and end states of these word HMMs are non-emitting)



Word HMMs



Grammar




Full HMM

# Questions from last time: Continuous speech recognition with HMMs

What is this grammar?

- A representation of all allowed sequences of words in the language (and optionally their probabilities)
- Allowed sequences are those that can be produced by traversing edges from the start state to the end state (in this case, these are the same single state)
- In this case, all transitions (edges) are equally probable



# Questions from last time: How does state tying affect EM training?

Recall: The Baum-Welch algorithm

- E step: Compute the *expected counts*

$$\sum_{t=1}^T \gamma_t(i) = \text{expected count of state } i$$

$$\sum_{t=1}^{T-1} \gamma_t(i) = \text{expected count of transitions from state } i$$

$$\sum_{t=1}^{T-1} \xi_t(i, j) = \text{expected count of transitions from } i \text{ to } j$$

# Questions from last time: How does state tying affect EM training?

Recall: The Baum-Welch algorithm M step

$$\hat{a}_{ij} = \frac{\text{expected-count}(i \rightarrow j)}{\text{expected-count}(\text{transitions from } i)} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}$$

$$\hat{b}_i(k) = \frac{\text{expected-count}(v_k \text{ in state } i)}{\text{expected-count}(i)} = \frac{\sum_{t=1, o_t=v_k}^T \gamma_t(i)}{\sum_{t=1}^T \gamma_t(i)}$$

$$\hat{\pi}_i = \text{expected-count}(q_1 = i) = \gamma_1(i)$$

## Questions from last time: How does state tying affect EM training?

E step (computation of expected counts)

- Counts are computed together for all tied states, e.g. if states  $i$  and  $j$  are tied:

$$\frac{1}{2} \left( \sum_{t=1}^T \gamma_t(i) + \sum_{t=1}^T \gamma_t(j) \right) = \text{expected count of state } i \text{ or } j$$

M step

- Unaffected!

## Questions from last time: How do we represent the state transition matrix for very large models?

- The state transition matrix is typically very sparse
- Usually, only transitions allowed are to self with probability  $p$  or next with probability  $1 - p$
- So, we need only store each state's self-transition probability

## Questions from last time: Where does supervised training come in?

HMM training with Baum-Welch is unsupervised, so how do we use labeled training sets?

Training an ASR system is more than just training a single HMM

- Typically, we have one HMM per word or phone label
- Supervised training: Spoken utterances with corresponding label sequences
- Start/end times of each label may or may not be given
- If start/end times are given, then the problem becomes one of training each of the word/phone HMMs on the collection of segments corresponding to that label
- If start/end times are not given:
  - Estimate the start and end times (do *forced alignment* – more on this later)
  - Consider these to be additional latent variables in the EM algorithm (rarely done)



## Recap: Continuous ASR with the Viterbi algorithm

Recognition = finding the most probable word string  $\mathbf{w}^*$  given the acoustic observations

$$\mathbf{O}: \mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmax}} p(\mathbf{w}|\mathbf{O})$$

From Bayes' rule,  $p(\mathbf{w}|\mathbf{O}) = \frac{p(\mathbf{O}|\mathbf{w})p(\mathbf{w})}{p(\mathbf{O})}$ . Therefore,

$$\begin{aligned}\mathbf{w}^* &= \underset{\mathbf{w}}{\operatorname{argmax}} p(\mathbf{w}|\mathbf{O}) \\ &= \underset{\mathbf{w}}{\operatorname{argmax}} p(\mathbf{O}|\mathbf{w})p(\mathbf{w})\end{aligned}$$

- The “whole-sentence” HMM gives  $p(\mathbf{O}|\mathbf{w})$ , the **acoustic model**
- $p(\mathbf{w})$  is the **language model**

## Recap: Continuous ASR with the Viterbi algorithm

Summing over all possible state sequences  $\mathbf{q}$  for the word string  $\mathbf{w}$ :

$$\begin{aligned}\mathbf{w}^* &= \operatorname{argmax}_{\mathbf{w}} p(\mathbf{O}|\mathbf{w})p(\mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{w}} \sum_{\mathbf{q}} p(\mathbf{O}|\mathbf{q}, \mathbf{w})p(\mathbf{q}|\mathbf{w})p(\mathbf{w})\end{aligned}$$

- $p(\mathbf{O}|\mathbf{q}, \mathbf{w})$  is given by the observation (emission) distribution
- $p(\mathbf{q}|\mathbf{w})$  is given by the state transition probabilities
- *Viterbi approximation*: Assume there is a single most probable state sequence  $\mathbf{q}^*$  such that all other  $\mathbf{q}$  contribute a negligible amount to the sum. Then

$$\mathbf{w}^* = \operatorname{argmax}_{\mathbf{w}} p(\mathbf{O}|\mathbf{q}^*, \mathbf{w})p(\mathbf{q}^*|\mathbf{w})p(\mathbf{w})$$

- So we can maximize jointly over  $\mathbf{w}$  and  $\mathbf{q}$ :

$$\mathbf{w}^*, \mathbf{q}^* = \operatorname{argmax}_{\mathbf{w}, \mathbf{q}} p(\mathbf{O}|\mathbf{q}, \mathbf{w})p(\mathbf{q}|\mathbf{w})p(\mathbf{w})$$

## Recap: Continuous ASR with the Viterbi algorithm

- To do continuous speech recognition, we can string together word HMMs according to a *grammar* or *language model*
- (Will get back to language modeling later)
- We will find the best state sequence using the Viterbi algorithm, and output the corresponding word string

## Questions from last time: Isn't it unnatural to chop up speech into distinct segments?

- As we've discussed, speech doesn't have sharp divisions between one word and the next, or one phone and the next
- Transitions from one sound/word to the next are often gradual
- When we run Viterbi, we make a decision about when each state/HMM start and end
- This is not great! Some solutions:
  - Some older work modified Viterbi to sum over multiple possible paths with slightly different start/end times – much slower, and not too much gained in performance (so the “distinct segment assumption” wasn't hurting HMMs too badly)
  - Newer neural models allow for fuzzier (or no) decisions about start and end times, e.g. neural encoder-decoders and connectionist temporal classification (CTC)
  - We'll discuss the latter in the coming weeks

# The need for subword units

Whole-word HMMs have some problems

- Cannot model unseen words, or even words seen too few times in training data
- Number of parameters is proportional to the vocabulary size
- (Note: The more parameters, the more data needed; rule of thumb:  $\sim 10$  data points per parameter)

⇒ Whole-word models mainly restricted to small-vocabulary tasks

## Subword units

- In the same way that sentences can be composed of whole-word HMMs, words can be composed of sub-word HMMs
- Units are then shared among words

Type of units	Approximate # (in English)
words	>100,000
phones	50
diphones	2,000
triphones	10,000
syllables	5,000

- Number of parameters now proportional to the number of sub-word units, not number of words

# Baseforms


- Each word can be represented as a sequence of phonemes, the word's *baseform*  
dogs → d ao g z
- Some words may need more than one baseform  
the → dh ah, dh iy  
either → iy dh er, ay dh er
- Each word's baseform(s) can be looked up in a dictionary
- All words in training and test data must be in the dictionary, or else we get them wrong
- Typically, baseform dictionaries are written by hand → prone to errors, inconsistencies, disagreements among linguists, ...

# The need for context-dependent units

- Recall: phonemes have variants (*allophones*)
  - Aspirated and un-aspirated stops: *pin* vs. *spin*
  - Allophones of /t/: *too* vs. *stew* vs. *butter* vs. *but* (the last with “unreleased” /t/)
- Phonemes are also influenced by surrounding context due to *co-articulation*, e.g. due to articulatory inertia
  - *keep*, *geese* (front /k/, /g/) vs. *coop*, *goose* (back /k/, /g/)




# The need for context-dependent units: Example



“keep”

# The need for context-dependent units: Example (2)



“coop”

## Triphones: Context-dependent phone models

- Model each phoneme in the context of its left and right neighbor:  $k-iy+p$  triphone is  $/iy/$  in the context of a preceding  $/k/$  and following  $/p/$  (not a model of the *sequence*  $/k\ iy\ p/$ )
- Build separate HMM for each triphone
- $\sim 50$  phonemes  $\Rightarrow 50^3$  triphones!
- Not all triphones are possible (e.g.,  $k-l+p$  is not)
- Still,  $\sim 10,000$  are possible and some are very rare
- Solution: tying!

## A comparison of model sizes

- Consider continuous-density HMMs where state observation distribution is Gaussian mixture model (GMM)
- Assume 10 diagonal Gaussians per state and 39-dimensional MFCC vectors
- Context-independent (CI) phone models
  - $\sim 50$  phonemes with  $\sim 3$  states each  $\Rightarrow \sim 150$  GMMs
  - 10 Gaussians per GMM  $\Rightarrow \sim 1500$  Gaussians
  - 39 dimensions per Gaussian, each requiring a mean and variance  $\Rightarrow \sim 120,000$  parameters
  - Manageable with  $\sim 3$  hours of speech training data
- Context-dependent (CD) triphone models
  - $\sim 10,000$  triphones with  $\sim 3$  states each  $\Rightarrow \sim 30,000$  GMMs  $\Rightarrow \sim 300,000$  Gaussians  $\Rightarrow \sim 24,000,000$  parameters  $\Rightarrow \sim 600$  hours of speech?
  - Actually... since some of the triphones are extremely rare, this may still not be enough!


## Tying via agglomerative (bottom-up) clustering

- Cluster triphones (or individual triphone states) into “sufficiently similar” groups
- Start with each triphone/state in cluster by itself, then iterate:
  - Find “closest” pair of clusters
  - Merge them into single cluster
- Until some convergence criterion
  - Distance between closest clusters is above some threshold
  - Number of data points in each cluster is above some threshold
- Different clusterings based on distance measure:
  - $a$  and  $b$  are individual triphones,  $A$  and  $B$  are clusters
  - $\mathcal{D}(a, b)$  is some distance function between triphones
  - *Single-linkage*:  $\mathcal{D}(A, B) = \min_{a \in A, b \in B} \mathcal{D}(a, b)$
  - *Complete-linkage*:  $\mathcal{D}(A, B) = \max_{a \in A, b \in B} \mathcal{D}(a, b)$

## Bottom-up vs. top-down clustering

- Agglomerative clustering is limited to the case where we have *some* examples of each item
- In the case of triphones, many may be unseen!
- Agglomerative clustering therefore tends not to be used for tying triphones
- But still used for other tasks, e.g. automatic discovery of speech units
- Alternative: top-down clustering, via decision trees

# Decision tree triphone clustering



## Decision tree construction

- Start with all frames for a given phone in a single node (the root)
- Find the best “question” for partitioning the data at a given node into two classes
- Repeat for each node
- Stop when there is insufficient data at each node, or when the best question isn’t helpful enough



## Details: What questions shall we ask?

- Typically involve identities of previous/following phones
- Or other attributes (e.g. *phonetic features*) of previous/following phones: voicing, nasality, place of articulation, manner of articulation, etc.
- In principle, can ask about any subset of the features that are relevant at the current node
  - Is the following phone in  $\{aa, k, l\}$  and the previous phone voiced, fricated, or nasal?
  - That's a lot of questions!  $(\sum_j (\#features_j))$
- In practice, use smaller set of pre-determined questions
  - Create in advance a list of “reasonable” questions
  - In principle, questions (and therefore tree splits) can be  $n$ -ary; in practice, binary

## When to stop?

- Cross-validation: Measure likelihood with different tree sizes on a held-out data set, choose the tree that maximizes likelihood on held-out data
- In practice, simple heuristics are often used:
  - Data at node has fewer than threshold  $T$  samples
  - Best question does not improve likelihood significantly (note: best question should *a/ways* improve likelihood somewhat!)

## Where does the data come from?

- In general, we don't have phonetic labels for each frame of training data
- One solution: Automatically align the data given some “simple” initial model, e.g. monophone-based recognizer
- Optionally, iterate: Once we have a better triphone model, re-align training data and re-build DT

## Other issues in sub-word modeling

### Pronunciation modeling

- Pronunciations often don't match the dictionary (dialects, non-native accents, conversational speech)
- New words (e.g. names) – how to get their baseforms?
- One idea: Apply letter-to-sound DTs
- More modern idea: Learn a neural “grapheme-to-phoneme” (G2P) model

### Language-specific issues beyond English

- Some languages have more predictable mapping from orthography to sounds: Spanish, French, German, Korean, ...  
⇒ pronouncing dictionaries virtually unnecessary!
- Some languages have *less* regular mapping from orthography to sounds, e.g. Mandarin Chinese
- ... or even an ambiguous mapping, e.g. Arabic, Hebrew

# Hybrid generative/discriminative models

- Generative models are models of the data generation process
$$p(\mathbf{O}, \mathbf{w}) = p(\mathbf{O}|\mathbf{w})p(\mathbf{w})$$
- Recognition is “inversion” of generation,
$$\mathbf{w}^* = \operatorname{argmax}_{\mathbf{w}} p(\mathbf{w}|\mathbf{O}) = \operatorname{argmax}_{\mathbf{w}} p(\mathbf{O}|\mathbf{w})p(\mathbf{w})$$
- Discriminative approaches attempt to solve the task more directly, e.g. by modeling  $p(\mathbf{w}|\mathbf{O})$  directly or even minimizing the intended error rate directly
- Main motivation for **hybrid** models:
  - Discriminative models are good! Let’s use them!
  - Discriminative models for *sequences* are hard! (More later)
  - Let’s combine discriminative **frame classifiers** with generative **sequence models**

# Hybrid generative/discriminative models


Typical approach:

- 1 Train a frame-based discriminative classifier of sub-word units (e.g. phones, phone states, triphone states) given some labeled training data,  $c^* = f_c(\mathbf{o})$ , where  $c$  is the class and  $\mathbf{o}$  is a frame feature vector
- 2 The output is a posterior probability  $p(c|\mathbf{o})$
- 3 Convert  $p(c|\mathbf{o})$  to something like an observation model (a “likelihood”):  $p(\mathbf{o}|c) \propto \frac{p(c|\mathbf{o})}{p(c)}$
- 4 Use the result in place of the observation model in an HMM

## Hybrid generative/discriminative models

- Most popular type of frame classifier by far: neural network
- Long line of research on hybrid HMM/NN models by groups at ICSI Berkeley, IDIAP, and elsewhere (e.g., Bourlard and Morgan 1994)
- Motivations for using NNs:
  - For right choice of NNs and training criterion, the output gives an estimate of the class posterior probabilities  $p(c|\mathbf{o})$
  - NNs are “universal function approximators”: Whatever the optimal classifier function  $f_c(\mathbf{o})$  is, there is some neural network that approximates it arbitrarily well (under mild assumptions)
- Main other (distant) competitor for frame classification: Support vector machines

## Linear classifiers



Let's start with a binary linear classifier (change of notation  $\mathbf{o} \rightarrow \mathbf{x}$ ):

$$\begin{aligned} f(\mathbf{x}) &= 1, \quad \mathbf{w} \cdot \mathbf{x} > 0 \\ &= 0, \quad \text{otherwise} \end{aligned}$$

- Can add an extra dimension to the input  $\mathbf{x}$  which is always 1, to induce a “bias”



# Linear classifiers: Learning weights with the perceptron algorithm [Rosenblatt 1957]

- 1 Initialize the weights  $\mathbf{w}$
- 2 At each iteration  $t$ , retrieve an example input  $\mathbf{x}_j$  in a training set, with corresponding output  $y_j$ :
  - Compute the output  $f(\mathbf{x}_j)$  using the current weights
  - Update the weights: for all nodes  $0 \leq i \leq n$ ,

$$w_i(t+1) = w_i(t) + \alpha(y_j - f(\mathbf{x}_j))x_{j,i}$$

- 3 Repeat step 2 until the average error  $\frac{1}{t} \sum_{j=1}^t |\hat{y}_j - f(\mathbf{x}_j)|$  is less than some threshold, or until some maximum iteration number  $t$

This does gradient descent on the perceptron loss

$$\sum_j \max(0, -\mathbf{y}_j \mathbf{w} \cdot \mathbf{x}_j)$$

## Nonlinear classifiers


What if the classification boundary between classes is not linear?

Add nonlinearity:

$$\begin{aligned} f(\mathbf{x}) &= 1, \quad \sigma(\mathbf{w} \cdot \mathbf{x}) > \tau \\ &= 0, \quad \text{otherwise} \end{aligned}$$


Examples:

- Heaviside (step, threshold):  $\sigma(z) = 0, z < 0; 1, \text{otherwise}$
- Logistic sigmoid:  $\sigma(z) = \frac{1}{1+e^{-z}}$
- Hyperbolic tangent:  $\sigma(z) = \tanh(z)$



# Multilayer perceptrons (MLP)


What if the nonlinear perceptron is still not enough to represent the decision boundary? Add more layers:



- Each node  $i$  in each layer  $l$  now outputs  $y_i^l = \sigma(\mathbf{w}_i^l \cdot \mathbf{y}_i^{l-1} + b_i^l)$  (with bias  $b$  now explicit)
- Or writing each layer's output as a vector:  
 $\mathbf{y}^l = \sigma_l(\mathbf{W}_l \mathbf{y}^{l-1} + \mathbf{b}^l)$ , where  $\sigma$  is applied element-wise
- Final output:  $\mathbf{y} = \mathbf{y}^L$  for an  $L$ -layer network


## Example decision boundaries

Here the activation functions are all threshold functions:



## Multi-class outputs

We've also expanded from a single output node to multiple nodes, to allow for more than a single class



Typical activation function: Softmax  $y_i = \frac{\exp(z_i)}{\sum_{j=1}^n \exp(z_j)}$

where  $z_i = \mathbf{w}_i \cdot \mathbf{x} + b_i$  (Note: layer indexing dropped;  $\mathbf{x}$  refers to the input of the current layer)

## Feedforward neural networks

More generally, a feedforward neural network (NN, or DNN) is any vector function  $f(\mathbf{x})$  of a vector input  $\mathbf{x}$  that can be written as a composition of the layers we've defined:

$$\begin{aligned}f(\mathbf{x}) &= \mathbf{y}^L \\ \mathbf{y}^l &= \sigma_l(\mathbf{W}_l \mathbf{y}^{l-1} + \mathbf{b}^l) \\ \mathbf{y}^0 &= \mathbf{x}\end{aligned}$$

## Training neural networks

The parameters are learned to minimize some loss, or measure of badness of the outputs

- A NN is an MLP if it is trained with perceptron loss (though often MLP is used to refer to any feedforward NN)
- Other losses (per training example):
  - Squared loss:  $\mathcal{L}_{SE} = \sum_i (y_i - \hat{y}_i)^2$
  - Cross-entropy loss (log loss):  $\mathcal{L}_{CE} = \sum_i \hat{y}_i \log y_i$  (typical for multi-class classification)
- Total loss is the sum of the loss over all training examples

## Training neural networks (2)

To minimize loss  $\mathcal{L}$ , we typically use gradient descent:

$$w(t+1) = w(t) - \eta \frac{\partial \mathcal{L}}{\partial w}$$

- $w$  is any single weight
- $\eta$  is a user-defined learning rate (set, e.g., by tuning on held-out data or via some rule of thumb)
- $\mathcal{L}$  is computed for 1 training example, a training subset (minibatch), or the full training set (a batch)
- If a single example or a minibatch is used, then this is stochastic gradient descent
- Gradient can be computed via the chain rule; this is called backpropagation
- Fortunately, for typical losses, the gradients are similar and simple for all weights
- ... and we often don't have to compute them as there are excellent toolkits to do it for us



## When is it enough complexity?

- In theory\*, one hidden layer is sufficient to approximate any output function arbitrarily well
- In practice\*\*, multiple hidden layers can be *very* helpful
- \* In theory, theory is the same as practice; in practice, it is almost never the case
- \*\* “In practice” = it may be much easier to learn parameters with multiple hidden layers\*\*\*
- \*\*\* There is some theory to this practice too

## More activation functions

Examples:

- ReLU:  $\sigma(z) = \max(0, z)$
- Softsign:  $\sigma(z) = \frac{z}{1+|z|}$
- Softplus:  $\sigma(z) = \log(1 + e^z)$

