

# TTIC 31110

## Speech Technologies

May 5, 2020

# Announcements

- HW3 due Friday 5/8 7pm
- Tutorial 3 (yesterday) materials available online
- Coming up: Term project
  - End of this week: Materials and guidelines available
  - Week 7: Project proposals due
  - Week 9: Project updates
  - Finals week: Final project presentations & reports
- Feel free to discuss topic ideas with me/Ankita at office hours, via email
- Seek out partners for term project (see survey results, discuss on canvas, etc.)

## Questions from last week/year

We mainly discussed training HMMs for the case of discrete observations; what about continuous observations (e.g. MFCCs)?

- Discrete HMMs can be used with vector quantized features (see Tutorial 2, Lecture 6)
- In a *continuous-density* HMM, the discrete observation probabilities,  $b_i(k)$ , are replaced by continuous densities  $b_i(\mathbf{o})$
- The observation distributions are typically Gaussian or mixture

of Gaussians: 
$$b_i(\mathbf{o}) = \sum_{k=1}^K c_{ik} \mathcal{N}(\mathbf{o} | \mu_{ik}, \Sigma_{ik}), \quad 1 \leq i \leq N$$

- In the forward/backward/Viterbi algorithms: same algorithms, just replace  $b_i(k)$  by the value of the corresponding density  $b_i(\mathbf{o})$
- For EM training, the update equations look a bit different.

# Recall: The Baum-Welch re-estimation formulas

M step, with multiple observation sequences  $\mathbf{O}^1, \dots, \mathbf{O}^L$

$$\hat{a}_{ij} = \frac{\sum_{l=1}^L \sum_{t=1}^{T-1} \xi_t^l(i, j)}{\sum_{l=1}^L \sum_{t=1}^{T-1} \gamma_t^l(i)}$$

$$\hat{b}_i(k) = \frac{\sum_{l=1}^L \sum_{t=1, o_t=v_k}^T \gamma_t^l(i)}{\sum_{l=1}^L \sum_{t=1}^T \gamma_t^l(i)}$$

$$\hat{\pi}_i = \frac{1}{L} \sum_{l=1}^L \gamma_1^l(i)$$

# Baum-Welch for continuous-density HMMs

Single-Gaussian case:  $b_i(\mathbf{o}) = \mathcal{N}(\mathbf{o}|\mu_i, \Sigma_i)$ ,  $1 \leq i \leq N$

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)} = (\text{same as for discrete HMMs!})$$

$$\begin{aligned} \hat{\mu}_i &= \frac{1}{\sum_{t=1}^T \gamma_t(i)} \sum_{t=1}^T \gamma_t(i) \mathbf{o}_t \\ &= (\text{same as Gaussian mixture update for “component” } i) \end{aligned}$$

$$\hat{\Sigma}_i = \frac{1}{\sum_{t=1}^T \gamma_t(i)} \sum_{t=1}^T \gamma_t(i) (\mathbf{o}_t - \hat{\mu}_i)(\mathbf{o}_t - \hat{\mu}_i)^T$$

# Baum-Welch for continuous-density HMMs

Gaussian mixture case:

- This is often referred to as an HMM/GMM
- Now state and Gaussian component index are both latent variables
- Update equations now involve  $\gamma_t(i, k) =$  posterior probability of being in component  $k$  in state  $i$  at time  $t$
- See Rabiner tutorial for equations

## Questions from last week/year

How does it all fit together? How do I go from a pile of data to a speech recognizer?

# Meta-algorithm 1: Training a (whole-word) HMM/GMM-based speech recognizer

(1) Given:

- Training set of  $L$  utterances (acoustic features + corresponding word transcriptions)
- Hyperparameters: # states per word, # Gaussians per state, HMM “topology” (which transition probabilities are 0)
- Initial parameter values (guess)

(2) Repeat until convergence:

- E step: For each training utterance  $l$ , run forward and backward algorithms and compute the  $\xi$ s
- M step: Update parameters according to the Baum-Welch equations
- Check convergence (e.g., likelihood not higher than previous iteration by some amount  $\delta$ )



## Meta-algorithm 2: Training and tuning a (whole-word) HMM/GMM-based speech recognizer

(1) Given:

- Training set of  $L$  utterances (acoustic features + corresponding word transcriptions)
- Development (held out/tuning) set of  $D$  utterances
- Set of allowed hyperparameters: range of # states per word, range of # Gaussians per state

(2) For each allowed combination of hyperparameters:

- Train recognizer using meta-algorithm 1
- Record performance (error rate) on dev set

(3) Choose trained recognizer with best dev-set performance

(In practice, there are more efficient ways to tune hyperparameters than the for-loop above (the above is a “grid search”))

## Questions from last week/year

What about silence?

- Treat it like just another word in the vocabulary
- Often we have one “word” for utterance-initial/utterance-final silence, and one for short inter-word silences
- One twist: We usually don’t have silences marked in training data, so we allow for optionally skipping the silences

## Details: Measuring performance

- Most common measure: Word error rate (WER)
- $WER = (\# \text{ substitutions} + \# \text{ deletions} + \# \text{ insertions}) / (\# \text{ wds in reference script})$
- Example:

REF: The \* dogs are barking now

HYP: The uh smogs \* barking \*

I S D D

- $WER = (1+1+2)/5 = 80\%$
- Can be computed efficiently using dynamic programming (like DTW, Viterbi)
- Note: WER can be above 100%

## Questions from last time

How are HMMs used for speech technologies besides speech recognition?

- **Unsupervised** learning, e.g. discovering sound units in a low-resource language
  - Take a pile of speech without transcriptions
  - Train a single HMM on all of it
  - Each state is a “sound unit”
  - Look for short/long repeated sequences of states to discover phones/words
- Speech synthesis
  - Much like HMMs for speech recognition, but trained on a **single** speaker’s speech
  - Some care needed to ensure continuity across synthesized frames
  - Possibly replaced by neural methods


# Recall: Hybrid generative/discriminative models

Typical approach:

- 1 Train a frame-based discriminative classifier of sub-word units (e.g. phones, phone states, triphone states) given some labeled training data,  $c^* = f_c(\mathbf{o})$ , where  $c$  is the class and  $\mathbf{o}$  is a frame feature vector
- 2 The output is a posterior probability  $p(c|\mathbf{o})$
- 3 Convert  $p(c|\mathbf{o})$  to something like an observation model (a “likelihood”):  $p(\mathbf{o}|c) \propto \frac{p(c|\mathbf{o})}{p(c)}$
- 4 Use the result in place of the observation model in an HMM
- 5 Most popular type of frame classifier by far: neural network


## Recall: Feedforward neural networks

A feedforward neural network (NN, or DNN) is any vector function  $f(\mathbf{x})$  of a vector input  $\mathbf{x}$  that can be written as a composition of simple “layers”



- Each node  $i$  in each layer  $l$  outputs  $y_i^l = \sigma(\mathbf{w}_i^l \cdot \mathbf{y}_i^{l-1} + b_i^l)$
- Or writing each layer's output as a vector:  
 $\mathbf{y}^l = \sigma_l(\mathbf{W}_l \mathbf{y}^{l-1} + \mathbf{b}^l)$ , where  $\sigma$  is applied element-wise
- (Letting  $\mathbf{y}^0 = \mathbf{x}$ )
- Final output:  $f(\mathbf{x}) = \mathbf{y} = \mathbf{y}^L$  for an  $L$ -layer network

## Recall: Multi-class outputs



Typical activation function for final layer: Softmax

$$y_i = \frac{\exp(z_i)}{\sum_{j=1}^n \exp(z_j)}$$

where  $z_i = \mathbf{w}_i \cdot \mathbf{x} + b_i$  (Note: layer indexing dropped;  $\mathbf{x}$  refers to the input of the current layer)

- Outputs are  $\geq 0$  and sum to 1, so can be thought of as class posterior probabilities  $p(\text{class } i | \mathbf{x})$

## Recap: Training neural networks

The parameters are learned to minimize some loss, or measure of badness of the outputs

- A NN is an MLP if it is trained with perceptron loss (though often “MLP” is used to refer to any feedforward NN)
- For multi-class classification (softmax output layer activation function), most common loss is cross-entropy loss

$$\ell_{CE} = - \sum_c y_c \log f_c(\mathbf{x}), \text{ where}$$

$\mathbf{x}$  is input vector for one example in training set

$y_c = 1$  if ground-truth label  $= c$ , 0 otherwise

$f_c(\mathbf{x})$  is our estimate of  $p(c|\mathbf{x})$

- Cross-entropy loss also called *log loss*, because  $\ell_{CE} = - \log f_{c^*}(\mathbf{x})$  where  $c^*$  is the ground-truth label
- Total loss is the sum of the loss over all training examples



## Aside: A teeny bit of information theory

- If  $X$  is a discrete random variable taking one of  $N$  values with probabilities  $p_1, \dots, p_N$ , respectively, then the **entropy** of  $X$  is

$$H(X) = - \sum_{i=1}^N p_i \log_2 p_i$$

- This is the average number of bits needed to represent  $X$
- If the distribution of  $X$  is uniform, then  $H(X) = \log_2 N$
- A related term is **perplexity**  $PP_p(X) = 2^{H(X)}$
- If the distribution of  $X$  is uniform, then what is  $PP(X)$

## Aside: A teeny bit of information theory

- The **cross-entropy** of a model distribution  $q$  with respect to a true distribution  $p$  is

$$H(p, q) = - \sum_{i=1}^N p_i \log_2 q_i$$

- This is the average number of bits needed to represent  $X$  drawn from  $p$  using a code optimized for  $q$
- Going back to cross-entropy loss:  
$$\ell_{CE} = - \sum_c y_c \log f_c(\mathbf{x})$$
- This is the cross-entropy between the true distribution  $y_c$  and our estimate of it  $f_c(\mathbf{x})$
- $y_c$  happens to be a very simple distribution:  $y_c = 1$  if true label =  $c$ , 0 otherwise
- If we had some other ground-truth distribution (“soft” labels), could still use cross-entropy
- But then it would not be equivalent to log loss

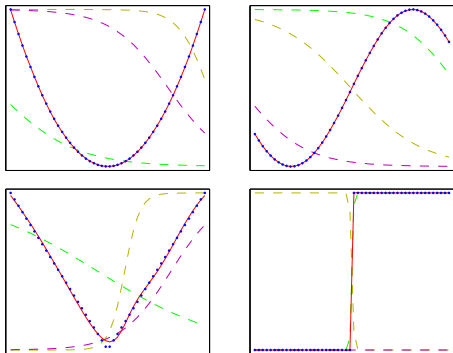
## Aside: A teeny bit of information theory

Cross-entropy loss:  $\ell_{CE} = -\sum_c y_c \log f_c(\mathbf{x})$

- Viewing the ground-truth label as a distribution over labels  $c$ , this is the cross-entropy between that distribution and the network's output distribution  $f_c(\mathbf{x})$
- This is a measure of dissimilarity between distributions
- (For our purposes, equivalent to KL divergence)
- What is the minimum of this loss? (Note if needed:  $x \log x \rightarrow 0$  as  $x \rightarrow 0$ )

## Power of two layers

- Theoretical result [Cybenko 1989]: 2-layer net with sigmoid hidden units can approximate any continuous function over compact domain to arbitrary accuracy, **given enough hidden units**
- Examples: 3 hidden units with  $\tanh(z) = \frac{e^{2z}-1}{e^{2z}+1}$  activation



[from Bishop]

# Back to speech recognition...

## Reminder: Hybrid ASR systems

- Use a DNN to produce a posterior for each class  $c$  (= HMM state) given an input frame of acoustic features  $\mathbf{o}$
- Posterior is converted to a scaled likelihood via  $p(\mathbf{o}|c) \propto \frac{p(c|\mathbf{o})}{p(c)}$

Where do class labels for all frames come from?

- Frames may have ground-truth (human) labels
- ... Or labels can be produced via a Viterbi alignment (“forced alignment”) using an existing HMM/GMM system
- ... Or we can use “soft labels” = posteriors produced by running forward-backward using an existing HMM/GMM system
- (The latter makes sense if using cross-entropy loss)

# Meta-algorithm 1a: Training a (whole-word) HMM/DNN-based speech recognizer

(1) Given:

- Training set of  $L$  utterances (acoustic features, corresponding word transcriptions, **state label per frame**)
- Hyperparameters: # states per word, # Gaussians per state, HMM “topology”, # **DNN layers**, # **DNN hidden units**, **learning rate**, **regularization parameters...**
- Initial parameter values for  $a_{ij}, \pi_i$  (**not**  $b_i(\mathbf{o})$ ),  $\Theta$

(2) **Train DNN: Repeat until convergence**

- One step of gradient descent
- Check convergence (e.g., loss not improved by at least  $\delta_{DNN}$ )

(3) **Train HMM: Repeat until convergence**

- E step: For each training utterance  $l$ , run forward and backward algorithms and compute the  $\xi$ s
- M step: Update parameters according to the Baum-Welch equations, **except**  $b_i(\mathbf{o})$


# Tandem models

Never mind the whole posterior conversion business: Use the NN outputs as features!

- Then use standard HMM/GMMs with these features as inputs
- Idea developed at ICSI Berkeley (e.g., Hermansky et al. 2000)
- $\mathbf{o}' = [y_1(\mathbf{o}) \ y_2(\mathbf{o}) \ \dots y_n(\mathbf{o})]$
- If the  $y_i(\mathbf{o})$  represent probabilities, then we typically take their logs:  $\mathbf{o}' = [\log(f_1(\mathbf{o})) \ \log(f_2(\mathbf{o})) \ \dots]$

## Tandem models (2)

Alternatively, use outputs from a lower layer, and make that layer narrow (a “bottleneck layer”) to reduce dimensionality



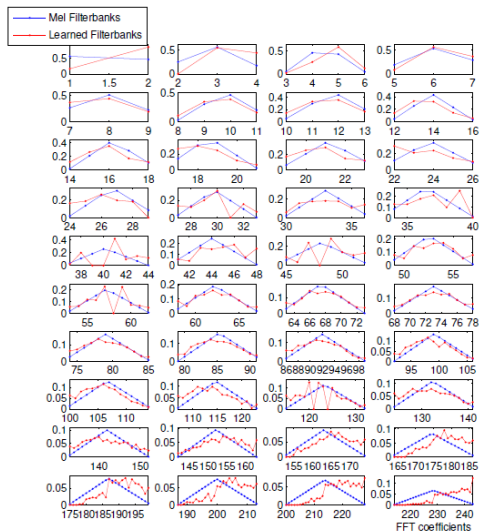


## Tandem models: More tricks (3)

- These features are often appended to the original features, e.g. MFCCs, so the new feature vector is  $[\mathbf{o} \ \mathbf{o}']$  (hence, “tandem”!)
- Typically, the input is a concatenation of acoustic vectors over a window of 7-20 frames around the current frame (very high-dimensional!)


# Visualizing learned acoustic features

Given raw spectrum input (Sainath *et al.* ASRU 2013):



# Visualizing learned acoustic features

Given a mel-spectrogram patch as input:



# Current state of hybrid and tandem models

As of 7-8 years ago:

- Depending on the task, HMM/NN models may or may not outperform HMM/GMM-based models
- Tandem models typically outperform their HMM/GMM-based counterparts

Now:


- DNNs have caused a revolution in ASR
- The state of the art is now often\* hybrid HMM/NN systems, with DNN-based tandem systems somewhat behind
- Tandem models have some advantages, e.g. easier to adapt to new speakers
- \*And for some domains, end-to-end neural network models are now the state-of-the-art

# What changed?

- More data
- More compute (GPUs)
- $\implies$  deeper networks
- $\implies$  wider network outputs
- Pretraining (that's probably not important, but was useful in getting NNs into the mainstream)
- Better regularization

## Recall: Visualizing learned acoustic features


Given a mel-spectrogram patch as input:



# Convolutional neural networks (CNNs)

But if we start from spectrograms, then it might help to consider:

- Many of the useful patterns are **local** in time-frequency
- Many of the useful patterns **repeat** in different time-frequency locations
- But they don't **exactly** repeat... they move around a bit between speakers, contexts, etc.




# Convolutional neural networks (CNNs)

CNNs are DNNs with some twists to take into account these considerations

- They encode **local** patterns with subsets of nodes that only consider small patches of the input (**filters**)
- They encode **repetition** of patterns by applying the same weights to different patches of the input (weight **sharing**)
- They normalize for **inexact repetition** by **pooling** information over multiple areas in the input
- Developed in the mid-1990s (really even the 1980s...) by Yann LeCun and colleagues
- Became hugely popular for image processing/computer vision starting in 2012
- Borrowed into speech recognition shortly thereafter




# CNNs



Key ingredients: local filters, sharing for repetition, pooling for inexact repetition


# CNNs

Often easier to think about convolutional layers in 2D




# CNNs


Convolutions can be applied to different kinds of features at the same time:



## CNNs: Example filters



# Dependence on number/types of layers



# of convolutional vs. fully connected layers	WER
No conv, 6 full (DNN)	21.6
1 conv, 5 full	21.3
2 conv, 4 full	18.9
3 conv, 3 full	20.2

# Dependence on activation type

**Table 12**

WER on broadcast news, 50 hr.

Model	Feature	Non-linearity	dev04f
GMM/HMM	fBMMI		18.8
DNN	fMLLR	sigmoid	16.3
CNN	log-mel	sigmoid	15.8
CNN+DNN	log-mel+(fMLLR+i-vectors)	sigmoid	14.2
CNN+DNN	log-mel+(fMLLR+i-vectors)	ReLU	<b>13.6</b>
DNN	log-mel+(fMLLR+i-vectors)	ReLU	14.2

**Table 13**

WER on broadcast news, 400 hr.

Model	Feature	Non-linearity	dev04f
GMM/HMM	fBMMI		16.0
DNN	fMLLR	sigmoid	15.1
CNN	log-mel	sigmoid	13.5
CNN+DNN	log-mel+(fMLLR+i-vectors)	ReLU	<b>12.7</b>

# CNNs applied to MFCCs

